Luis Gimenez

Data scientist · forecasting, causal inference and LLM applications

Kraków, Poland · lgimenezuso@gmail.com · https://www.linkedin.com/in/luis-gimenez-uso · https://github.com/luisgius

Last updated 23 Oct 2026

I'm a data scientist at Uber in Kraków. I work on driver supply forecasting, and on turning business questions into hypotheses that causal inference and experimental design can test.

Before Uber I was a junior data scientist at Buynomics in Cologne, building revenue prediction and clustering models for a Tier-1 CPG client. Before that I interned at Porsche AG in Stuttgart, automating test-bench data collection and building Power BI dashboards, and at Capgemini in Valencia as an ERP consultant. I studied computer science and business administration as a dual degree at the Universidad Politécnica de Valencia.

I also like to write about philosophy.

Experience

Data scientist at Uber, previously at Buynomics, with internships at Porsche AG and Capgemini, specializing in machine learning forecasting, causal inference and production-grade LLM applications.

Based in	Kraków, Poland
Time zone	Central European Time (UTC+1; UTC+2 in summer)
Work rights	European Union: EU citizen
Languages	English (full professional proficiency, C1) German (beginner, A2) Spanish (native) Polish (beginner)

2026 – now

Data Scientist

Uber, Kraków, Poland

Forecasting driver supply, and using causal inference and experimental design to guide supply-side operational decisions.

- Built and deployed an ML forecasting system for driver supply (LightGBM quantile regression with conformal prediction intervals) at city × product × week granularity, outperforming the incumbent manual forecasting process.
- Formulated business questions as testable hypotheses, applying causal inference and experimental design to estimate effects and guide supply-side operational decisions.

- Engineered features and cleaned large-scale datasets (Python, SQL over Presto and Spark), and shipped model overrides through a shadow-to-live rollout, validating performance against production before launch.
- Visualized and reported findings to stakeholders, delivering forecast reliability analyses and full project documentation for technical and business audiences.

Stack: Python, SQL, LightGBM, Presto, Spark

2025 – 2026

Junior Data Scientist

Buynomics, Cologne, Germany

Revenue prediction and clustering models for a Tier-1 CPG client, and the data pipelines that feed the Buynomics ML platform.

- Built and deployed revenue prediction and clustering models (Python, PyTorch) for a Tier-1 CPG client, including a logistic purchase-decision component, mapping transaction data for 500+ SKUs.
- Designed Python ETL workflows to connect raw country-level data to the Buynomics ML platform, improving data reliability and model training speed.
- Defined statistical evaluation metrics to benchmark model accuracy, ensuring production-ready outputs for commercial decision-making.

Stack: Python, PyTorch

2024 – 2025

Data Analyst – BI & Visualization (internship)

Porsche AG, Stuttgart, Germany

Internship automating test-bench data collection and building data models and Power BI dashboards for R&D and QA.

- Developed Python scripts to automate test bench data collection, reducing manual processing time by ~70% (from hours to minutes) and ensuring consistent data quality.
- Built predictive KPI models, designing data models and Power BI dashboards that give the R&D and QA departments real-time performance visibility.

Stack: Python, Power BI

2023 – 2024

ERP Consultant (internship)

Capgemini, Valencia, Spain

Internship as an ERP consultant for an enterprise healthcare client.

- Translated business requirements into technical specifications and designed cross-departmental workflows for an enterprise healthcare client; the documentation was adopted as an organizational standard.

Projects

Pragmatic Bayes and Bounded Frequentism · 2026
On a simulated counting experiment whose signal cannot be negative, a neural posterior came close to the exact posterior's Bayes risk in under a microsecond per observation.
https://github.com/luisgius/bayes-frequentist

Stock market RL lab · 2024 – now
A PPO agent choosing between an ETF and cash lost to buy-and-hold after trading costs, and a pre-declared five-ETF replication came back adverse.
https://github.com/luisgius/stock-market-rl-lab

AI-Powered RAG Recommendation System · 2024
My Bachelor's thesis, a retrieval-augmented recommender that explains its suggestions, with lexical, vector and hybrid search compared on ranking metrics and an LLM-as-judge.

Education

2018 – 2024
B.Sc. Computer Science and B.Sc. Business Administration (dual degree)
Universidad Politécnica de Valencia, Valencia

Double major: two full bachelor's degrees completed concurrently (300+ ECTS), combining software engineering with business strategy. Erasmus exchange in International Business Management at UBB, Cluj-Napoca, Romania.

Certifications

Machine Learning Specialization · Stanford Online / DeepLearning.AI

Python for Data Science · Maven Analytics

Financial Markets · Yale University

Skills

Machine Learning & Statistics

Python (Pandas, NumPy, Scikit-learn), LightGBM, PyTorch, SQL, Time-Series Forecasting, Quantile Regression, Logistic Regression (GLM), Conformal Prediction, Causal Inference, Experimental Design, A/B Testing, Predictive Modeling, Clustering

Generative AI & LLMs

RAG Architectures, LangChain, LangGraph, OpenAI/Anthropic APIs, Vector Databases (FAISS), Hybrid Search (BM25 + Embeddings), Retrieval & Relevance Evaluation (NDCG, MRR, Precision@k), Prompt Engineering

ML Engineering & Cloud

Distributed Data Processing (Spark, Presto), Docker, FastAPI, Git, AWS, Azure, API Deployment, ETL Pipelines, Data Visualization, Power BI